

Relational Programming in SQL

PL/pgSQL - SQL Procedural Language

Dirk Van Gucht¹

¹Indiana University

October 26, 2018

PL/pgSQL Relational Programming Environment

- Types
- Variables
 - Simple-type variables; int, text, boolean etc
 - Complex-type variables; composite, array, JSON, etc
 - Relation and view variables
 - Cursors
- Functions and Triggers
- Declaration statements
 - Type declarations (CREATE TYPE)
 - Simple and complex-type variable declarations (DECLARE)
 - Relation and view variable declaration (CREATE TABLE and CREATE VIEW)
 - Functions and triggers declarations (CREATE FUNCTION and CREATE TRIGGER)
 - Cursor declaration (DECLARE CURSOR)

Garbage collection statements

- Garbage collection for types (DROP TYPE)
- Garbage collection for relation and view variables (DROP TABLE, DROP VIEW)
- Garbage collection for functions (DROP FUNCTION)
- Garbage collection for triggers (DROP TRIGGER)
- Garbage collection for cursors (DROP CURSOR)

Control statements

- Return statement (RETURN *expressions*;))
- PERFORM query
- Block statements (BEGIN . . . END)
- Loop statements (LOOP, WHILE, FOR)
- Conditional statements (IF THEN ELSE, CASE)
- Cursor operations (OPEN, FETCH, CLOSE)

Relational Programs

The skeleton of a relational program is as follows:

```
CREATE OR REPLACE FUNCTION functionName (list of arguments)
  RETURNS return type AS
$$
(Optional Declaration section;)
BEGIN
  sequence of statements;
END;
$$ LANGUAGE plpgsql;
```

Example: Relational Program

An example with the **IF** statement

```
CREATE OR REPLACE FUNCTION convert(a char)
  RETURNS integer AS
$$
BEGIN
  IF (a = 't') THEN RETURN 1;
  ELSE
    IF (a = 'f') THEN RETURN 0;
    ELSE
      IF (a = 'u') THEN RETURN 1/2;
      ELSE RETURN(2);
    END IF;
  END IF;
END IF;
END;
$$ LANGUAGE plpgsql;
```

→

SELECT convert('u');
convert
1/2
SELECT convert('z');
convert
2

Example: Relational Program

An example with the **CASE** statement

```
CREATE OR REPLACE FUNCTION convert(a char)
  RETURNS integer AS
$$
BEGIN
  CASE WHEN (a = 't') THEN RETURN 1;
        WHEN (a = 'f') THEN RETURN 0;
        WHEN (a = 'u') THEN RETURN 1/2;
        ELSE RETURN 2;
  END CASE;
END;
$$ LANGUAGE plpgsql;
```

→

SELECT convert('u');
convert
1/2
SELECT convert('z');
convert
2

Example: Iterative Program

Iterative program for the **factorial(n)** function

```
CREATE OR REPLACE FUNCTION factorial_iterative (n integer)
RETURNS integer AS
$$
DECLARE
    result integer;
    i integer;
BEGIN
    result := 1;
    FOR i IN 1..n
    LOOP
        result := i * result;
    END LOOP;
    RETURN result;
END;
$$ language plpgsql;
```


Example: Iterative Program

Recursive program for the **factorial(n)** function

```
CREATE OR REPLACE FUNCTION factorial_Recursive (n integer)
RETURNS integer AS
$$
BEGIN
    IF n = 0 THEN
        RETURN 1;
    ELSE
        RETURN n * factorial_Recursive(n-1);
    END IF;
END;
$$ language plpgsql;
```

Functions that affect the database state

- Functions can be defined to affect (change) the database state
- Often such functions do not need to return values: they have **VOID** return type

```
CREATE OR REPLACE FUNCTION change_db_state()  
  RETURNS VOID AS  
$$  
BEGIN  
  DROP TABLE foo_relation;  
  CREATE TABLE foo_relation(a integer);  
  INSERT INTO foo_relation VALUES (1), (2), (3);  
  DELETE FROM foo_relation WHERE a=1;  
END;  
$$ language plpgsql;
```

```
→      select change_db_state();  
      change_db_state  
-----  
      select * from foo_relation;  
      a  
-----  
      2  
      3
```

Example: Local functions

- You can also **CREATE local functions**
- Care must be taken careful with function **delimiters**

```
CREATE OR REPLACE FUNCTION globalFunction()  
  RETURNS void AS  
$proc$  
BEGIN  
  CREATE OR REPLACE FUNCTION localFunction()  
    RETURN integer AS  
  $$  
    SELECT 5;  
  $$ language sql;  
END; $proc$ language plpgsql;
```

→

```
SELECT globalFunction();  
      globalfunction
```

```
SELECT localFunction();  
      localfunction  
      5
```

Two kinds of assignment statements

- The typical assignment statement is of the form
 $x := \text{expression};$
- An assignment to a variable can also be done with a query and the clause

$\text{SELECT INTO variable query};$

- The value of the query is assigned to the variable

```
CREATE OR REPLACE FUNCTION size_of_A FUNCTION foo ()  
  RETURNS integer AS  
$$  
  DECLARE counter integer;  
  BEGIN  
    SELECT INTO counter COUNT(*) from A;  
    RETURN counter;  
  END;  
$$ language plpgsql
```

→

SELECT * FROM A;
x

'A'
'B'
SELECT size_of_A();
size_of_A

2

SELECT INTO (non-deterministic behavior)

- SELECT INTO can lead to **non-deterministic (random)** effects!
- This is because SELECT INTO **chooses the first available tuple from the result of the query** and assigns it to the INTO variable (in our case the variable `element_from_A`).¹
- Of course, this can be useful when **sampling** data

```
CREATE OR REPLACE FUNCTION choose_one_from_A()  
  RETURNS integer AS  
$$  
  DECLARE element_from_A integer;  
  BEGIN  
    SELECT INTO element_from_A a.x  
    FROM (SELECT x from A ORDER BY random()) a;  
    RETURN element_from_A;  
  END;  
$$ language plpgsql
```

→

```
SELECT choose_one_from_A();  
choose_one_from_a  
-----  
'B'  
  
SELECT choose_one_from_A();  
choose_one_from_a  
-----  
'A'  
  
SELECT choose_one_from_A();  
choose_one_from_a  
-----  
'A'
```

¹ If the query does not return any tuple, then the variable is set to NULL.

"Assignment" statements to relation variables

- "Assignment" statements to relation (table) variables are done using the **INSERT INTO**, **DELETE FROM**, and **UPDATE** statements

```
CREATE OR REPLACE FUNCTION relation_assignment()
  RETURNS void AS
$$
BEGIN
  CREATE TABLE IF NOT EXISTS AB(A integer, B integer);
  DELETE FROM AB;
  INSERT INTO AB VALUES (0,0);
  INSERT INTO AB SELECT a1.x, a2.x FROM A a1, A a2;
  UPDATE AB SET A = A*A WHERE B = 2;
END;
$$ language plpgsql;
```

→

```
select * from A;
  x
-----
  1
  2

SELECT * FROM AB;
ERROR: relation "ab" does not exist

SELECT relation_assignment();
SELECT * from AB;
  a  b
----
  0  0
  1  1
  2  1
  1  2
  4  2
```

Iterators over collections

- Relations and arrays are **collections**
- Relations are unordered collections whereas arrays are ordered collections
- We consider **iterator variables** that slide (move; iterate) over such a collection **one element at a time**
- In SQL, an iterator variable over a relation (which may or may not be the result of a query) is often referred to as a **CURSOR**
- In SQL, it is frequently not necessary to use cursors as the following function illustrates

```
CREATE OR REPLACE FUNCTION there_is_book_that_cost_more_than(k integer)
  RETURNS boolean AS
$$
BEGIN
  SELECT EXISTS(SELECT * FROM book WHERE price > k);
END;
$$ language sql
```

Iterators over collections (**cursors**)

```
CREATE OR REPLACE FUNCTION there_is_book_that_cost_more_than(k integer)
  RETURNS boolean AS
$$
BEGIN
  SELECT EXISTS(SELECT * FROM book WHERE price > k);
END;
$$ language sql;
```

The following function with the same semantics does use the
iterator record variable (cursor) **b**

```
CREATE OR REPLACE FUNCTION there_is_book_that_cost_more_than(k integer)
  RETURNS boolean AS
$$
DECLARE exists_book boolean;
        b RECORD;
BEGIN
  exists_book := false;
  FOR b IN SELECT * FROM book
  LOOP
    IF b.price > k
    THEN exists_book := true;
    EXIT;
    END IF;
  END LOOP;
  RETURN exists_book;
END; $$ language plpgsql;
```


Iterators over arrays

- Below is an example from the PostgreSQL manual illustrating iteration through an array using the **FOR EACH** clause
- The function **sum** takes an integer array as input and returns the sum of its elements
- The variable **x** is the iterator which gets assigned, one at a time, to each **element** in the array
- Note in particular that **x** is **not assigned to index positions** of the array

```
CREATE FUNCTION sum(A int[])
  RETURNS int8 AS
  $$ DECLARE
    s int8 := 0;
    x int;
  BEGIN
    FOR EACH x IN ARRAY A
    LOOP
      s := s + x;
    END LOOP; RETURN s;
  END;
  $$ LANGUAGE plpgsql;
```

Iterators over arrays

On the right is an alternative version for the `sum` function. There an index variable is used that iterates over the index positions of the array.

```
CREATE FUNCTION sum(A int[])
  RETURNS int8 AS
$$ DECLARE
  s int8 := 0;
  x int;
BEGIN
  FOR EACH x IN ARRAY A
  LOOP
    s := s + x;
  END LOOP;
RETURN s;
END;
$$ LANGUAGE plpgsql;
```

⇔

```
CREATE FUNCTION sum(A int[])
  RETURNS int8 AS
$$ DECLARE
  s int8 := 0;
  i int;
BEGIN
  FOR i IN array_lower(A,1)..array_length(A,1)
  LOOP
    s := s + A[i];
  END LOOP;
RETURN s;
END;
$$ LANGUAGE plpgsql;
```

The Ancestor-Descendant Relation

- Assume that we are given a **parent-child relation**

$PC(\text{parent}, \text{child})$

In this relation, a pair (p, c) indicates that person p is the parent of person c

- Starting from the information in the PC relation, we want to compute the ancestor-descendant

$ANC(\text{ancestor}, \text{descendant})$

relation.

- In this relation, a pair (a, d) indicates that person a is an ancestor of person d .

Computing the Ancestor-Descendant relation in Pure SQL

- It can be shown that there does not exist any Pure SQL query that can compute the ANC relation
- However, if we add an iteration construct to Pure SQL then this becomes possible

A recursive definition of the ANC relation

- The ANC relation can be recursively defined using the following rules:

Rule 1: if $\text{PC}(p, c)$ then $\text{ANC}(p, c)$

Rule 2: if $\text{ANC}(a, p)$ and $\text{PC}(p, c)$ then $\text{ANC}(a, c)$

- Rule 1 states that if p is a parent of c then p is an ancestor of c
- Rule 2 states that if a is an ancestor of p and if p is the parent of c , then a is an ancestor of c

Computing the ANC relation in stages

Rule 1: if $\text{PC}(p, c)$ then $\text{ANC}(p, c)$

Rule 2: if $\text{ANC}(a, p)$ and $\text{PC}(p, c)$ then $\text{ANC}(a, c)$

- These two rules allow us to compute the ANC relation in stages:
 - Stage 1: Start with the (parent,child) pairs in PC
 - Stage 2: Add to these the (grandparent,grandchild) pairs
 - Stage 3: Add to these the (great-grandparent, great-grandchild) pairs
 - etc
- This computation in stages will terminate because the PC relation is assumed to be a finite relation
- We will show how we can compute these stages using a simple iterative process

Computing the ANC relation in stages with RA expressions

- We will specify the computation at each stage using an RA expression
- We assume that PC has attributes P and C and that for each i , ANC_i has attributes A and D

$$\begin{aligned}ANC_0 &\leftarrow PC \\ANC_1 &\leftarrow ANC_0 \cup \pi_{A,C}(ANC_0 \bowtie_{D=P} PC) \\ANC_2 &\leftarrow ANC_1 \cup \pi_{A,C}(ANC_1 \bowtie_{D=P} PC) \\&\dots \\ANC_{i+1} &\leftarrow ANC_i \cup \pi_{A,C}(ANC_i \bowtie_{D=P} PC) \\&\dots \\&\text{until } ANC_{n+1} = ANC_n\end{aligned}$$

- Clearly, this sequence of computation suggest a simple loop

Computing ANC incrementally

- To develop our code, we first specify a function

`new_ANC_pairs()`

- This function computes, given an approximation for the ancestor relation ANC , additional (ancestor, descendent)-pairs that should be present in ANC
- Such pairs can be discovered by **joining** (i.e., composing) the current approximation for ANC with the PC relation

Computing ANC incrementally

- More precisely, at stage i ($i \geq 1$), the function `new_ANC_pairs()` computes $(\text{ANC}_i - \text{ANC}_{i-1})$
- Observe that to compute ANC_i , it then suffices to insert into ANC_{i-1} this new set of new pairs
- Further observe that when the function `new_ANC_pairs()` at some stage $n+1$ returns no new pairs, then $\text{ANC}_{n+1} = \text{ANC}_n$ and, in fact, then $\text{ANC}_n = \text{ANC}$

The `new_ANC_pairs()` function in SQL

Recall that `ANC` has schema (A, D) and `PC` has schema (P, C)

```
CREATE OR REPLACE FUNCTION new_ANC_pairs()
  RETURNS TABLE (A integer, D integer) AS
$$
  (SELECT A, C
   FROM   ANC, PC
   WHERE D = P)
EXCEPT
  (SELECT A, D
   FROM   ANC);
$$ language sql;
```

The Ancestor_Descendant () program

We can now specify a function (a program) that computes the ANC relation

```
CREATE OR REPLACE FUNCTION Ancestor_Descendant()  
  RETURNS void AS  
$$  
BEGIN  
  DROP TABLE IF EXISTS ANC;  
  CREATE TABLE ANC(A integer, D integer);  
  
  INSERT INTO ANC SELECT * FROM PC;  
  
  WHILE EXISTS (SELECT * FROM new_ANC_pairs())  
  LOOP  
    INSERT INTO ANC SELECT * FROM new_ANC_pairs();  
  END LOOP;  
END;  
$$ language plpgsql;
```

Alternative version using SQL Recursive View

- Observe that the `ANC` relation can also be computed using a recursive view as follows:

```
WITH RECURSIVE ANC(A,D) AS  
  (SELECT P, C FROM PC  
   UNION  
   SELECT A, C  
   FROM   ANC INNER JOIN PC ON (D = P) )  
  
SELECT A, D FROM ANC;
```

Illustration

PC	
P	C
1	2
1	3
1	4
2	5
2	6
3	7
5	8

SELECT Ancestor_Descendant();
→
SELECT * FROM ANC;

ANC	
A	D
1	2
1	3
1	4
1	5
1	6
1	7
1	8
2	5
2	6
2	8
3	7
5	8

Paths and Connectivity in Graphs

- A **graph** is a pair $G = (V, E)$ where V is a set of **vertices** (nodes) V and E is a set of **edges**
- The set of edges E is a subset of $V \times V$. Consequently, $|E| \leq |V|^2$
- A **path** in G is a sequence of its vertices (v_0, v_1, \dots, v_n) such that
 - 1 $n \geq 1$, and
 - 2 for each $i \in [1, n - 1]$, (v_i, v_{i+1}) is an edge in E
- n is called the **length** of the path.

Paths and Connectivity in Graphs

- Let G be a graph let s and t be a pair of its vertices.
- We say that s and t are **connected** in G if there exists a path (v_0, v_1, \dots, v_n) in G such that $s = v_0$ and $t = v_n$.
- We say that G is **connected** if each of its pairs of vertices is connected.

Path Connectivity in a Graph

- Assume that we are given a graph $G = (V, E)$ where E is represented by the relation `Edge(source, target)`
- In this relation, a pair (s, t) indicates that $(s, t) \in E$.
- We want to compute the connectedness relation

`Connected(source, target)`

- In this relation, a pair (s, t) indicates that there is a path in G from vertex s to vertex t

Path Connectivity in a Graph

- We can define the `(Connected)` recursively using the following two rules:

Rule 1: if `Edge(s, t)` then `connected(s, t)`

Rule 2: if `connected(s, u)` and `Edge(u, t)`
then `connected(s, t)`

- Rule 1 states that if there is an edge from s to t then s is connected to t
- Rule 2 states that if s is connected to a vertex u and if there is an edge from u to t , then s is connected to t

Path Connectivity in a Graph

- Notice that these are exactly the same rules as those for the Ancestor-Descendant relation!
- So if we replace in the previous program the $PC(P, C)$ relation by $Edge(source, target)$ and the relation $ANC(A, D)$ by the relation $Connected(source, target)$, then we have program to compute the set of all pairs of vertices in G that are connected
- The relation $Connected$ is often called the **transitive closure** of G and is denoted by TC .²

²We will assume that the scheme of TC is $(source, target)$.

Path Connectivity in a Graph (Transitive Closure)

```
CREATE OR REPLACE FUNCTION new_TC_pairs()
  RETURNS TABLE (source integer, target integer)AS
$$
BEGIN
  (SELECT TC.source, Edge.target
   FROM   TC INNER JOIN Edge ON (TC.target = Edge.source)
  EXCEPT
  (SELECT source, target
   FROM TC);
END;
$$ language plpgsql;
```

Path Connectivity in a Graph (Transitive Closure)

We can now specify a function that computes the TC relation

```
CREATE OR REPLACE FUNCTION TC()
  RETURNS void AS
$$
BEGIN
  DROP TABLE IF EXISTS TC;
  CREATE TABLE TC(source integer, target integer);

  INSERT INTO TC SELECT * FROM Edge;

  WHILE EXISTS (SELECT * FROM new_TC_pairs())
  LOOP
    INSERT INTO TC SELECT * FROM new_TC_pairs();
  END LOOP;
END;
$$ language plpgsql;
```

Illustration

Edge	
source	target
1	2
1	3
2	3
3	2
3	4
3	5
6	7
6	6
7	8
4	9
9	5

Initialization of TC (Phase 1)
Paths of length 1
→

TC	
source	target
1	2
1	3
2	3
3	2
3	4
3	5
6	7
6	6
7	8
4	9
9	5

Illustration

length 1 paths
TC

source	target
1	2
1	3
2	3
3	2
3	4
3	5
6	7
6	6
7	8
4	9
9	5

Phase 2
new pairs
→

length 2 paths
new_TC_pairs

source	target
1	4
1	5
2	2
2	4
2	5
3	3
3	9
4	5
6	8

Phase 2
UNION
→

TC

source	target
1	2
1	3
1	4
1	5
2	2
2	3
2	4
2	5
3	2
3	3
3	4
3	5
3	9
4	5
4	9
6	6
6	7
6	8
7	8
9	5

Illustration

length 1 and 2 paths

TC	
source	target
1	2
1	3
1	4
1	5
2	2
2	3
2	4
2	5
3	2
3	3
3	4
3	5
3	9
4	5
4	9
6	6
6	7
6	8
7	8
9	5

Phase 3
new pairs
→

length 3 paths

new_TC_pairs	
source	target
1	9
2	9

Phase 3
UNION
→

TC	
source	target
1	2
1	3
1	4
1	5
1	9
2	2
2	3
2	4
2	5
2	9
3	2
3	3
3	4
3	5
3	9
4	5
4	9
6	6
6	7
6	8
7	8
9	5

Illustration

Edge	
source	target
1	2
1	3
2	3
3	2
3	4
3	5
6	7
6	6
7	8
4	9
9	5

SELECT TC();
→
SELECT * FROM TC;

TC	
source	target
1	2
1	3
1	4
1	5
1	9
2	2
2	3
2	4
2	5
2	9
3	2
3	3
3	4
3	5
3	9
4	5
4	9
6	6
6	7
6	8
7	8
9	5

Transitive Closure (non-linear recursion)

- Observe that `connected` can also be specified by the following two rules:

Rule 1 if `Edge(s, t)` then `connected(s, t)`

Rule 2 if `connected(s, u)` and `connected(u, t)`
 then `connected(s, t)`

- Rule 1 states that if there is an edge in E from s to t then s and t are connected.
- Rule 2 states that if s is connected to a vertex u and if u is connected to t , then s is connected to t .

Transitive Closure (non-linear recursion)

- The nice aspect of this recursive definition is that $O(\log(n))$ applications of these rules suffice to compute the transitive closure (i.e., the connectedness relation).
- Here n denotes the length of the longest path in the graph
- $O(n)$ number of applications may be required when we use the original rules for defining the connectedness relation

Path Connectivity in a Graph (Transitive Closure) Non-linear recursion

Notice how the `new_TC_Pairs` function uses non-linear recursion the `FROM TC p1 INNER JOIN TC p2` clause

```
CREATE OR REPLACE FUNCTION new_TC_pairs()  
  RETURNS TABLE (source integer, target integer) AS  
$$  
BEGIN  
  (SELECT p1.source, p2.target  
   FROM   TC p1 INNER JOIN TC p2 ON (p1.target=p2.source)  
  EXCEPT  
  (SELECT source, target  
   FROM   TC);  
END;  
$$ language plpgsql;
```

Path Connectivity in a Graph (Transitive Closure) Non-linear version

The TC function remains the same

```
CREATE OR REPLACE FUNCTION TC()
  RETURNS void AS
$$
BEGIN
  DROP TABLE IF EXISTS TC;
  CREATE TABLE TC(source integer, target integer);

  INSERT INTO TC SELECT * FROM Edge;

  WHILE EXISTS (SELECT * FROM new_TC_pairs())
  LOOP
    INSERT INTO TC SELECT * FROM new_TC_pairs();
  END LOOP;
END;
$$ language plpgsql;
```