

MapReduce simulation in SQL of the word_count problem

Dirk Van Gucht

1 The Word Count problem

- A **document** is modeled as a $(doc, words)$ pair
 1. doc the document identifier of a document
 2. $words$ is the set (bag) of words in that document
- The input to the **word count** is a set of documents
- The output is a set of pairs $(word, wordCount)$ where $wordCount$ is the number of occurrences of the word $word$ across all the documents
- We will present a SQL program in **MapReduce**-style for the **word count** problem.

- The relation `documents` is created as follows. Notice that we represent a bag of words with an array

```
CREATE TABLE documents(
  doc text,
  words text[]);
```

- Populate the `documents` relation as follows¹

```
INSERT INTO documents VALUES('d1', ARRAY['A','B','C']);
INSERT INTO documents VALUES('d2', ARRAY['B','C','D']);
INSERT INTO documents VALUES('d3', ARRAY['A','E']);
INSERT INTO documents VALUES('d4', ARRAY['B','B','A','D']);
INSERT INTO documents VALUES('d5', ARRAY['E','F']);
```

We get

doc		words
-----+-----		
d1		{A,B,C}
d2		{B,C,D}
d3		{A,E}
d4		{B,B,A,D}
d5		{E,F}

¹Notice that a word may occur multiple times in a document.

- The expected output for the word count problem is the following:

word		word_count
-----+-----		
A		3
B		4
C		2
D		2
E		2
F		1

2 The Word Count problem in SQL

- Before we show the MapReduce simulation, we begin by writing a SQL program for the [word count](#) problem
- This program will serve as a blueprint for the MapReduce simulation.

```
WITH
%map:      doc_word AS (SELECT word, 1 AS one
                                FROM   documents d, UNNEST(d.words)),
%group:    word_ones AS (SELECT word, array_agg(one) AS ones
                                FROM   doc_word
                                GROUP BY (word),
%reduce:    occurrences AS (SELECT word, CARDINALITY(ones) AS wordCount
                                FROM word_ones p)

% output:
SELECT    * FROM occurrences
```

3 MapReduce programs

- A **basic** MapReduce program is a pair of functions (**mapper**, **reducer**).
- The **mapper** function takes as input a $(key, value)$ -pair and outputs a set (bag) of $(key, value)$ -pairs.
- Note that the output *key* and *values* need not be of the same type as the input *key* and *value*.
- The **reducer** function takes as input a (key, bag_of_values) -pair and outputs a set (bag) of $(key, value)$ -pairs.

4 Semantics of a MapReduce program

- The semantics of a **basic** MapReduce program consists of a **map-**, a **group-**, and a **reduce-phase**:
 - In the **map-phase**, the mapper is **map-applied**² to a set of $(key, value)$ pairs that come from the input and the outputs of all these calls are put together in a binary relation $A(key, value)$.
 - In the **group-phase**, the A relation is grouped³ on its *key*-value column⁴, and for each *key*-value, a pair (key, bag_of_values) is produced, where the *bag_of_values* is the bag of all values in A with that *key* value.
 - In the **reduce-phase**, the reducer is map-applied⁵ to the (key, bag_of_values) pairs produced in the group-phase. For each such pair, the reducer produces a bag of $(key, value)$ pairs.
 - The **output** of the program is the bag of all these $(key, value)$ pairs.
- The semantics of a MapReduce program is the **composition**⁶ of basic MapReduce programs.

²Typically in a parallelized manner.

³Sometimes also called *shuffled*

⁴Typically by hashing on this key-value.

⁵Typically in a parallelized manner.

⁶Typically in a pipe-lined manner.

5 Simulating a (basic) MapReduce program in SQL

- In Section 2, we expressed the **word count** problem in SQL. We deliberately wrote it in a fashion that resembles the map-, group-, and reduce- phases present in the semantics of a basic MapReduce program.
- We can give an even more faithful simulation if we write this SQL program using a mapper function and a reducer function. This can be done with SQL user-defined functions.

- We specify the **mapper** function as follows:

```
CREATE OR REPLACE FUNCTION mapper(doc text, words text[])
    RETURNS TABLE (word text, one integer) AS
$$
    SELECT w.wd, 1 FROM (SELECT UNNEST(words) AS wd) w;
$$ LANGUAGE SQL;
```

- The **mapper** function does the following when applied to a document:

```
SELECT p.word, p.one FROM mapper('d1',array['A','A','B']) p;
```

word		one
A		1
A		1
B		1

- The **mapper** function when map-applied to the **documents** relation produces the relation **map_output(word, one)**.

```
WITH map_output AS
    (SELECT word, one
     FROM documents d, LATERAL(SELECT p.word, p.one
                               FROM mapper(d.doc,d.words) p) q)

SELECT word, one
FROM map_output;
```

- Notice how we use **LATERAL** clause. This is necessary since in the **FROM** clause we need to apply the **mapper** function to each document **d** in the **documents** relation (i.e., the input pairs)

The map_output relation consist of the following key value pairs:

word		one
-----+-----		
A		1
B		1
C		1
B		1
C		1
D		1
A		1
E		1
B		1
B		1
A		1
D		1
E		1
F		1

- Before we specify the **reducer** function, we show how the **group-phase** prepares the inputs to this function.
- This will be done by taking the **map_output** relation and grouping it on **word**
- This will associate with each **word** the bag of 1-values that it occurs with in this relation.
- These bags of 1-values will be formed using the **ARRAY** aggregation operator.
- We will put the output of the group-phase in the relation **group_output(word,ones)**.

```
WITH group_output AS
    (SELECT p.word, array_agg(p.one)
     FROM   map_output p
     GROUP BY (p.word))
```

```
SELECT word, ones FROM group_output;
```

word	ones
F	{1}
A	{1,1,1}
E	{1,1}
C	{1,1}
B	{1,1,1,1}
D	{1,1}

- We now specify the **reducer** function. In our case, this function takes as input a (**word**, bag of 1's **ones**) pair and outputs the desired (word, wordCount) pair
- For ease of programming, we let the reducer function return a relation with this (**word**, **wordCount**) pair.

```
CREATE OR REPLACE FUNCTION reducer(word TEXT, ones INTEGER[])
RETURNS TABLE(word TEXT, word_count INTEGER) AS
$$
    SELECT reducer.word, CARDINALITY(ones);
$$ LANGUAGE SQL;
```

- The **reducer** function does the following when applied to a (**word**, **ones**) pair.

```
SELECT word, word_count
FROM   reducer('A', '{1,1,1,1}');
```

word	word_count
A	4

- We can now map-apply the **reducer** function to the (word, ones) pairs **b** generated in the **group-by phase** and get the desired output.

```
SELECT t.word, t.word_count
FROM   group_output r, LATERAL(SELECT s.word, s.word_count
                                FROM   reducer(r.word, r.ones) s) t
```

word		word_count
F		1
A		3
E		2
C		2
B		4
D		2

6 MapReduce simulation in SQL

- Putting everything together we get the following SQL simulation of the **word_count** MapReduce program.

```
WITH
  %mapper phase
  map_output AS
  (SELECT q.word, q.one
   FROM   documents d, LATERAL(SELECT p.word, p.one
                                FROM   mapper(d.doc,d.words) p) q),

  %group phase
  group_output AS
  (SELECT p.word, array_agg(p.one)
   FROM   map_output p
   GROUP BY (p.word),

  %reducer phase
  reduce_output AS
  (SELECT t.word, t.word_count
   FROM   group_output r, LATERAL(SELECT s.word, s.word_count
                                FROM   reducer(r.word, r.ones) s) t)

  %output
  SELECT word, word_count
  FROM   reduce_output;
```

7 MapReduce in distributed setting

- In a distributed setting of compute nodes connected by a network, the **documents** relation is stored in **chunks** across the local file systems of these nodes.
- A mapper can process the chunk of documents at its compute node and then send its output to other compute nodes. This is typically done by applying a hash-function to a key-value. This hash-function will give the location of another compute node. The (key-value) pair is then sent to the compute node with the key's hash-function value.
- After all the appropriate values for a key have been sent to the appropriate compute nodes, the reducers can go to work locally (at the compute node) on the list of values associated with a key.
- The reducers can transmit their output, or they can keep it locally for further processing by other MapReduce programs.
- A big problem is skew in the data. It is possible that there is an uneven distribution of the values associated with keys. In that case, computation can slow considerably and the benefits of parallel (distributed) computing can be lost.